

[ArXiv2020]Linformer: Self-Attention with Linear Complexity

Facebook AI, Seattle, WA

- Category
 - 벤다이어그램 : Linformer , Linera Transformer, Performer \in Low Rank / Kernels
 - 그래프 : Linformer, MCA \in Memory compress

1. Introduction

Model Architecture	Complexity per Layer	Sequential Operation
Recurrent	$O(n)$	$O(n)$
Transformer, (Vaswani et al., 2017)	$O(n^2)$	$O(1)$
Sparse Tansformer, (Child et al., 2019)	$O(n\sqrt{n})$	$O(1)$
Reformer, (Kitaev et al., 2020)	$O(n \log(n))$	$O(\log(n))$
Linformer	$O(n)$	$O(1)$

2. Related Work

- Transformer and self-attention

$$\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V) = \underbrace{\text{softmax} \left[\frac{QW_i^Q (KW_i^K)^T}{\sqrt{d_k}} \right]}_P VW_i^V$$

Sparse Attention (Child et al., 2019): This technique improves the efficiency of self-attention by adding sparsity in the context mapping matrix P . For example, the Sparse Transformer (Child et al., 2019) only computes P_{ij} around the diagonal of matrix P (instead of the all P_{ij}). Meanwhile, blockwise self-attention (Qiu et al., 2019) divides P into multiple blocks and only computes P_{ij} within the selected blocks. However, these techniques also suffer a large performance degradation, while having only limited additional speed-up, i.e., 2% drop with 20% speed up.

LSH Attention (Kitaev et al., 2020): Locally-sensitive hashing (LSH) attention utilizes a multi-round hashing scheme when computing dot-product attention, which in theory reduces the self-attention complexity to $O(n \log(n))$. However, in practice, their complexity term has a large constant 128^2 and it is only more efficient than the vanilla transformer when sequence length is extremely long.

3. Self-Attention is Low Rank

기본 세팅

- We use two pretrained transformer models, RoBERTa-base (12-layer stacked transformer) and RoBERTa-large (24-layer stacked transformer) on two tasks
- two tasks: masked-language-modeling task on Wiki103 and classification task on IMDB
- ([RoBERTa: A Robustly Optimized BERT Pretraining Approach](#))

Empirical view




Figure 1: Left two figures are spectrum analysis of the self-attention matrix in pretrained transformer model (Liu et al., 2019) with $n = 512$. The Y-axis is the normalized cumulative singular value of context mapping matrix P , and the X-axis the index of largest eigenvalue. The results are based on both RoBERTa-base and large model in two public datasets: Wiki103 and IMDB. The right figure plots the heatmap of normalized cumulative eigenvalue at the 128-th largest eigenvalue across different layers and heads in Wiki103 data.

[Theroem 1] self-attention is low rank

Theorem 1. (self-attention is low rank) For any $Q, K, V \in \mathbb{R}^{n \times d}$ and $W_i^Q, W_i^K, W_i^V \in \mathbb{R}^{d \times d}$, for any column vector $w \in \mathbb{R}^n$ of matrix VW_i^V , there exists a low-rank matrix $\tilde{P} \in \mathbb{R}^{n \times n}$ such that

$$\Pr(\|\tilde{P}w^T - Pw^T\| < \epsilon\|Pw^T\|) > 1 - o(1) \text{ and } \text{rank}(\tilde{P}) = \Theta(\log(n)), \quad (3)$$

where the context mapping matrix P is defined in (2).

- Proof

$$P = \text{softmax} \underbrace{\left[\frac{QW_i^Q (KW_i^K)^T}{\sqrt{d}} \right]}_A = \exp(A) \cdot D_A^{-1},$$

▼ Johnson–Lindenstrauss lemma (JL lemma) - from stanford cs369m lecture1.pdf

3 Johnson-Lindenstrauss Lemma

Lemma For any $0 < \epsilon < 1$ and any interger n let k be a possitive interger such that

$$k \geq \frac{24}{3\epsilon^2 - 2\epsilon^3} \log n \quad (2)$$

then for any set A of n points $\in \mathbb{R}^d$ there exists a map $f : \mathbb{R}^d \rightarrow \mathbb{R}^k$ such that for all $x_i, x_j \in A$

$$(1 - \epsilon)\|x_i - x_j\|^2 \leq \|f(x_i) - f(x_j)\|^2 \leq (1 + \epsilon)\|x_i - x_j\|^2 \quad (3)$$

◦ $\tilde{P} = PR^T R$

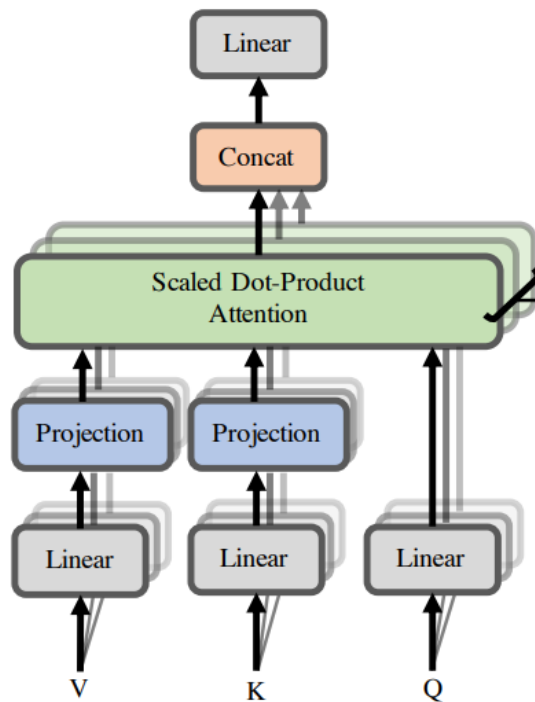
- Let R be an $k \times n$ matrix, $1 \leq k \leq n$, with i.i.d. entries from $N(0, 1/k)$


-
- **Given the low-rank property of the context mapping matrix P** , one straightforward idea is to use singular value decomposition (SVD) to approximate P with a low-rank matrix P_{low}

$$P \approx P_{low} = \sum_{i=1}^k \sigma_i u_i v_i^T = \underbrace{\begin{bmatrix} u_1, \dots, u_k \end{bmatrix}}_U \text{diag}\{\sigma_1, \dots, \sigma_k\} \underbrace{\begin{bmatrix} v_1 \\ \vdots \\ v_k \end{bmatrix}}_V \quad (6)$$

- One can use P_{low} to approximate self-attention (2) with error and $O(nk)$ time and space complexity.
- However, this approach **requires performing an SVD decomposition in each self-attention matrix**, which adds additional complexity.
 - ▼ SVD time complexity
 - SVD : $O(m^2n + mn^2 + n^3)$
 - truncated SVD : $O(mn^2)$
 - reference : <https://mathoverflow.net/questions/161252/what-is-the-time-complexity-of-truncated-svd>
- Therefore, we propose **another approach for low-rank approximation** that avoids this added complexity.

4. Model





- Add two linear projection matrices $E_i, F_i \in \mathbb{R}^{n \times k}$ when computing key and value : $(n \times d) \rightarrow (k \times d)$

$$\begin{aligned} \overline{\text{head}}_i &= \text{Attention}(QW_i^Q, E_iKW_i^K, F_iVW_i^V) \\ &= \underbrace{\text{softmax}\left(\frac{QW_i^Q(E_iKW_i^K)^T}{\sqrt{d_k}}\right)}_{\bar{P}: n \times k} \cdot \underbrace{F_iVW_i^V}_{k \times d} \end{aligned}$$

- only require $O(nk)$ time and space complexity
- if $k \ll n$, then we can significantly reduce the memory and space consumption.
- The following theorem states that, when $k = O(d/2)$ (independent of n), one can approximate $P \cdot VW_i^V$ using linear self-attention with error.

[Theorem 2] Linear self-attention

Theorem 2. (Linear self-attention) For any $Q_i, K_i, V_i \in \mathbb{R}^{n \times d}$ and $W_i^Q, W_i^K, W_i^V \in \mathbb{R}^{d \times d}$, if $k = \min\{\Theta(9d \log(d)/\epsilon^2), 5\Theta(\log(n)/\epsilon^2)\}$, then there exists matrices $E_i, F_i \in \mathbb{R}^{n \times k}$ such that, for any row vector w of matrix $QW_i^Q(KW_i^K)^T/\sqrt{d}$, we have

$$\Pr(\|\text{softmax}(wE_i^T)F_iVW_i^V - \text{softmax}(w)VW_i^V\| \leq \epsilon \|\text{softmax}(w)\| \|VW_i^V\|) > 1 - o(1) \quad (8)$$

- Proof

Proof. Define $E = \delta R$ and $F = e^{-\delta} R$, where $R \in \mathbb{R}^{n \times k}$ with i.i.d. entries from $N(0, 1/k)$, δ is a constant with $\delta = 1/2^n$. We will first prove that for any row vector $x \in \mathbb{R}^n$ of matrix QK^T and column vector $y \in \mathbb{R}^n$ of matrix V ,

$$\Pr (\| \exp(xE^T) F y^T - \exp(x) y^T \| \leq \epsilon \| \exp(x) y^T \|) > 1 - 2e^{-(\epsilon^2 - \epsilon^3)k/4}. \quad (21)$$

- Parameter sharing between projections

(1) Headwise sharing : $E_i = E$ and $F_i = F$ across all heads i

(2) Key-value sharing : $E_i = F_i = E$ across all heads i

(3) Layerwise sharing : E

ex) in a 12-layer, 12-head stacked Transformer model, headwise sharing, key-value sharing and layerwise sharing will introduce 24, 12, and 1 distinct linear projection matrices, respectively.

- Nonuniform projected dimension

- As shown in figure 1(right), heads in higher layer tends towards a more skewed distributed spectrum (lower rank).
- This implies one can choose a smaller projected dimension k for higher layers.

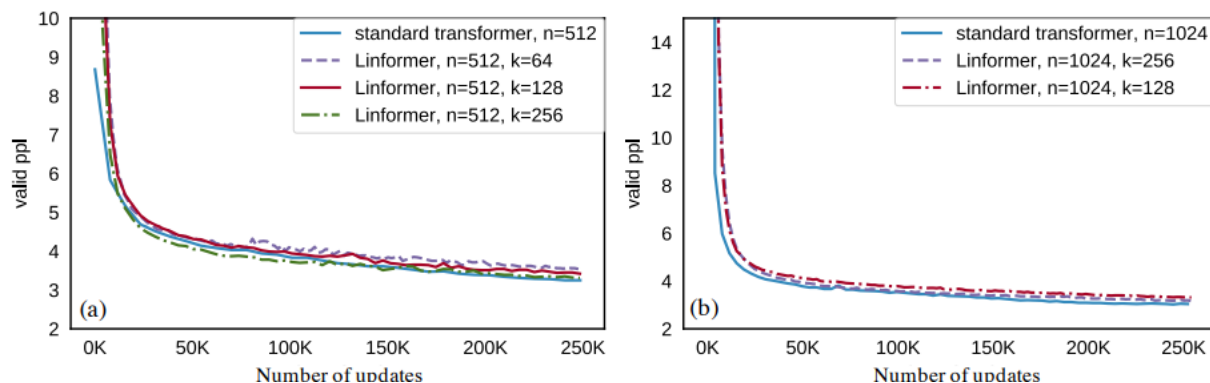
5. Experiments

5.1 Pretraining Perplexities

- Compare the pretraining performance of our proposed architecture against RoBERTa based on the Transformer
- we use BookCorpus + English Wikipedia as our pretraining set (3300M words).
- All models are pretrained with the masked-language-modeling (MLM) objective
- parallelized across 64 Tesla V100 GPUs with 250k updates.

(1) Effect of projected dimension


- different k & $n=512/1024$



- As expected, the Linformer performs better as projected dimension k increases
- However, even at $k = 128$ for $n = 512$ and $k = 256$ for $n = 1024$, Linformer's performance is already nearly on par with the original Transformer

(2) Effect of Sharing projections

- $n=512$

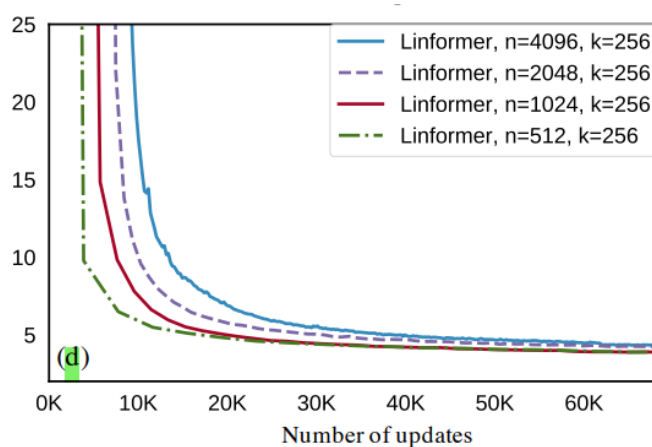


- when we use just a single projection matrix (i.e. for layerwise sharing), the resulting Linformer model's validation perplexity almost matches that of the non-shared model

→ we can decrease the number of additional parameters in our model, and consequently, it's memory consumption, without much detriment to performance.

(3) Effect of longer sequences

- $n \in \{512, 1024, 2048, 4096\}$ & k fixed at 256.



- As sequence length increases, even though our projected dimension is fixed, the final perplexities after convergence remain about the same.
- This further empirically supports our assertion that the Linformer is linear-time.

5.2 Downstream Results

- We finetune our Linformer on
 - IMDB (sentiment classification)
 - SST-2 (sentiment classification)
 - QNLI (natural language inference)
 - QQP (textual similarity)

n	Model	SST-2	IMDB	QNLI	QQP	Average
512	Liu et al. (2019), RoBERTa-base	93.1	94.1	90.9	90.9	92.25
	Linformer, 128	92.4	94.0	90.4	90.2	91.75
	Linformer, 128, shared kv	93.4	93.4	90.3	90.3	91.85
	Linformer, 128, shared kv, layer	93.2	93.8	90.1	90.2	91.83
	Linformer, 256	93.2	94.0	90.6	90.5	92.08
	Linformer, 256, shared kv	93.3	93.6	90.6	90.6	92.03
	Linformer, 256, shared kv, layer	93.1	94.1	91.2	90.8	92.30
512	Devlin et al. (2019), BERT-base	92.7	93.5	91.8	89.6	91.90
	Sanh et al. (2019), Distilled BERT	91.3	92.8	89.2	88.5	90.45
1024	Linformer, 256	93.0	93.8	90.4	90.4	91.90
	Linformer, 256, shared kv	93.0	93.6	90.3	90.4	91.83
	Linformer, 256, shared kv, layer	93.2	94.2	90.8	90.5	92.18

Table 2: Dev set results on benchmark natural language understanding tasks. The RoBERTa-base model here is pretrained with same corpus as BERT.

- $n = 512, k = 128$: comparable to RoBERTa
- $n = 512, k = 256$: even slightly outperforms RoBERTa
 - although the Linformer’s layerwise sharing strategy shares a single projection matrix across the entire model, it actually exhibits the best accuracy
- $(n = 1024, k = 256) \sim (n = 512, k = 256)$
 - empirically supports that the performance of Linformer model is mainly **determined by the projected dimension k instead of the ratio n/k .**

5.3 Inference-time Efficiency Results

- benchmark both models’ inference speed and memory on a 16GB Tesla V100 GPU card.
- We randomly generate data up to some sequence length n and perform a full forward pass on a multiple batches.
- We also choose batch size based on the maximum batch size that can fit in memory

length n	projected dimensions k					length n	projected dimensions k				
	128	256	512	1024	2048		128	256	512	1024	2048
512	1.5x	1.3x	-	-	-	512	1.7x	1.5x	-	-	-
1024	1.7x	1.6x	1.3x	-	-	1024	3.0x	2.9x	1.8x	-	-
2048	2.6x	2.4x	2.1x	1.3x	-	2048	6.1x	5.6x	3.6x	2.0x	-
4096	3.4x	3.2x	2.8x	2.2x	1.3x	4096	14x	13x	8.3x	4.3x	2.3x
8192	5.5x	5.0x	4.4x	3.5x	2.1x	8192	28x	26x	17x	8.5x	4.5x
16384	8.6x	7.8x	7.0x	5.6x	3.3x	16384	56x	48x	32x	16x	8x
32768	13x	12x	11x	8.8x	5.0x	32768	56x	48x	36x	18x	16x
65536	20x	18x	16x	14x	7.9x	65536	60x	52x	40x	20x	18x

Table 3: Inference-time efficiency improvements of the Linformer over the Transformer, across various projected dimensions k and sequence lengths n . Left table shows time saved. Right table shows memory saved.

- We also plot inference times of both Linformer and Transformer on the 100 data samples in the top right of Figure 2

